

Lecture 5

Local Search

Thanapon Noraset
Faculty of ICT, Mahidol University

Adapted from AIMA by Stuart Russell and Peter Norvig, and UC Berkeley CS188 by Dan Klein and Pieter Abbeel (ai.berkeley.edu)

Agenda





- Local Search
 - Hill Climbing
 - Random restart
 - Local Beam Search
- Simulated Annealing
- Genetic Algorithms

Exploring the state space



- Search algorithms so far explore the state space systematically
- They also keep track of multiple alternative paths (i.e. frontiers, legal moves)



Local Search




Advantages of Local Search



- Work with problems where a solution is single state, not a path.
- Use very little memory (usually a constant).
- Often find reasonable solutions in a large or infinite state space (continuous).

From a current node

- Check local changes \rightarrow Neighbors in state space
- Select a better neighbor until you cannot find any better



Objective Functions




- Local search algorithms need to compare “scores” of different options
- **Objective function** or heuristic function provide these score
- The goal states have the maximum (or minimum) score
- For example
 - 8-Queens: number of attacking pairs
 - Eval function: average root mean square
(vs actual minimax values)

Example: N-Queens




- Start with all the queens
- Neighbor states = all possible States to move a single queen
 - How many for N queens?



- Heuristic function: the number of pairs of queens that are attacking each other

Example: N-Queens

How many pairs are attacking each other?



Hill Climbing Search



- General Idea:
 - Start at a random state
 - Keep moving to **the best neighbor**
 - If no better neighbors, exit
- Advantage:
 - Simple and memory efficient

function HILL-CLIMBING(*problem*) **returns** a state that is a local maximum

current \leftarrow MAKE-NODE(*problem*.INITIAL-STATE)

loop do

neighbor \leftarrow a highest-valued successor of *current*

if *neighbor*.VALUE \leq *current*.VALUE **then return** *current*.STATE

current \leftarrow *neighbor*

Example: N-Queens



18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14	♔	13	16	13	16
♔	14	17	15	♔	14	16	16
17	♔	16	18	15	♔	15	♔
18	14	♔	15	15	14	♔	16
14	14	13	17	12	14	12	18

Current cost = ?

Neighbor states
and their heuristic
costs

Best move cost = ?

Example: N-Queens



18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14	♔	13	16	13	16
♔	14	17	15	♔	14	16	16
17	♔	16	18	15	♔	15	♔
18	14	♔	15	15	14	♔	16
14	14	13	17	12	14	12	18

Current cost = 17

Nearby states
and their heuristic
costs

Best move cost = 12

Hill Climbing Search




- Not complete nor optimal
- Get stuck in local optimal states and plateaus
- Only solve 14% of 8-Queen starting states
- Very fast, on average it takes
 - 4 steps when solvable
 - 3 steps when stuck


Hill Climbing Search: Stuck




Every neighbor has higher cost



State-space Landscape





Where do you end up if you start from:

- X, Y, or Z

stop ↓ ↓ ↓
B D E

*B, E : local optimal - best among its neighbor
some are global*

Sideway: Keep moving on on a plateau

- Need to limit how many sideway steps
 - Why?
- 8-Queens (limit = 100 moves):
 - Improve from 14% to 94% success rate
 - Increase average steps to 21 / 64


Randomness comes to rescue!

- Stochastic choice:
Randomly choose from random uphill moves
- Random-restart:
“If at first you don’t succeed, try, and try again”
 - On average, it takes 7 restarts to solve 8-Queens

Local Beam Search *: bit complicated no need to do in hw*

- Instead of a single current state, we can keep track k best states at once
- At each step pick the best k neighbors (regardless of the branch)
- Commonly used in speech recognition and machine translation

Local Beam Search: Example



“Come over here, the grass is greener!”

Simulated Annealing



- Hill climbing search gets stuck because it *never* makes downhill moves
- Moving randomly is complete (if we wait long enough), it is too inefficient
- Simulated Annealing combines both
 - Escape local maxima by allowing downhill moves
 - But make less downhill moves as time goes on

Simulated Annealing: Algorithm



function SIMULATED-ANNEALING(*problem*, *schedule*) **returns** a solution state

inputs: *problem*, a problem

schedule, a mapping from time to “temperature”

current \leftarrow MAKE-NODE(*problem*.INITIAL-STATE)

for $t = 1$ **to** ∞ **do**

$T \leftarrow \text{schedule}(t)$

if $T = 0$ **then return** *current*

next \leftarrow a randomly selected successor of *current*

$\Delta E \leftarrow \text{next.VALUE} - \text{current.VALUE}$


if $\Delta E > 0$ **then** *current* \leftarrow *next*

else *current* \leftarrow *next* only with probability $e^{\Delta E/T}$


Simulated Annealing





- Lower temperature \Rightarrow Fewer “downhill” moves
- If T decreased slowly enough, it will converge to optimal state!
- ... theoretically ... In practice, it can stuck too :-)



Genetic Algorithm




- Sometimes it is helpful to move far away, rather than jittering around the same place.
- Genetic algorithm defines a way to “jump” around using “natural selection” analogy
- Need to encode a state into a sequence



83742516

Genetic Algorithm: Idea



- Start with a random population (fix size)
 - State → Individual (gene)
- States with higher **fitness function** will have higher chance to reproduce
- Reproduction by **crossover** and **mutation**

Genetic Algorithm



function GENETIC-ALGORITHM(*population*, FITNESS-FN) **returns** an individual

inputs: *population*, a set of individuals

FITNESS-FN, a function that measures the fitness of an individual

repeat

new_population \leftarrow empty set

for $i = 1$ **to** SIZE(*population*) **do**

$x \leftarrow$ RANDOM-SELECTION(*population*, FITNESS-FN)

$y \leftarrow$ RANDOM-SELECTION(*population*, FITNESS-FN)

child \leftarrow REPRODUCE(x, y)

if (small random probability) **then** *child* \leftarrow MUTATE(*child*)

add *child* to *new_population*

population \leftarrow *new_population*

until some individual is fit enough, or enough time has elapsed

return the best individual in *population*, according to FITNESS-FN


function REPRODUCE(x, y) **returns** an individual

inputs: x, y , parent individuals

$n \leftarrow$ LENGTH(x); $c \leftarrow$ random number from 1 to n

return APPEND(SUBSTRING($x, 1, c$), SUBSTRING($y, c + 1, n$))

Genetic Algorithm: N-Queens



- Similar to stochastic local beam search + crossover operation
- But sometimes crossover provides no advantage at all
 - Crossover only makes sense when we can make independent improvement on a block of the gene code

- Local search algorithms
 - Used when a solution is a state
 - Start from a random state and move to a neighbor state
- Algorithms:
 - Hill Climbing Search + its variants
 - Simulated Annealing
 - Genetic Algorithms

Quiz



See MyCourse